

CP Commands

***Please note, we are able to see multiple request from engineers related to performing some PUT and POST operation (curl put/post permission has been revoked by service team). However, we would like to inform that some of these command can be executed using ‘aes domain dp node es safe-curl post’ and ‘aes domain dp node es curl cluster settings’**

```
PUT setting works with "aes domain dp node es curl cluster settings":
--url "_cluster/settings" --data '{"transient" : {"indices.breaker.total.limit" :
"99%"} }'
--url "_cluster/settings" --data '{ "transient" :
{"aes.jetty.admission_control.global_jvmp.limit": "99%"} }'
--url '/_cluster/settings' --data '{"persistent" : {"cluster.blocks.read_only" :
false}}'

POST snapshot restore works with "aes domain dp node es safe-curl post":
--url "_snapshot/cs-automated-enc/<snapshot_id>/_restore" -D $DOM --data '{"indices":
"index_1,index_n"}'

# - low check di vs dm
aescloudopstools $REGION domain cp cs-domain-manager-status --args "-D $DI"
aescloudopstools $REGION domain cp cs-domain-manager-status --args "-D $DI -ia --all"

# - moderate retry failed activities on dm
aescloudopstools $REGION domain cp cs-retry-failed-activities --args "-D $DI"

# - low check vfi
aescloudopstools $REGION domain cp cs-list --args "-D $DI --instances -k rev"
aescloudopstools run $REGION domain cp cs-manage-software-release-config --args "--action list" | awk '{print $1,$2}'

# - low print snapshots
aescloudopstools $REGION domain cp cs-manage-es-snapshots-read-only --args "-D $DI"

# - low print shard allocation
aescloudopstools $REGION domain cp es-print-shards-info --args "-D $DI"

# - low print instances dataNodes
tumbler $REGION aes domain dp get-domain-resources -D $DI | \
jq -r '.ec2Instances[] | select(.nodeType=="dataNode") | .instanceId' | \
awk '{print $1}'

# - low print instances masterNodes
tumbler $REGION aes domain dp get-domain-resources -D $DI | \
jq -r '.ec2Instances[] | select(.nodeType=="masterNode") | .instanceId' | \
awk '{print $1}'

# - low max allowed di count
aescloudopstools $REGION generic run --cmd-name "cs-create-domain-instance" --cmd-args "-D $DI --max-allowed-di-
count 3"

# - high check and fix quorum loss
aescloudopstools run $REGION generic run --cmd-name detect-quorum-loss.py --cmd-args "-D $DI"
aescloudopstools $REGION generic run --cmd-name es-master-script.py --cmd-args "-D $DI -A FixQuorumLoss --force" --
timeout 3600

# - low broken list whitelist hourly snapshots
aescloudopstools $REGION domain cp cs-whitelist-client-snapshot --feature
es.snapshot.frequent_snapshot_disabled_domains \
--args "--target-stack swift-$REGION-prod --target-client-id -r prod --region $REGION --action list"
# - 2021/04/28 - Commands above for list and disable hourly snapshots is broken since python3
```

```
# - workaround is to mwinit --breakglass -o and connect in ssh into a Ressource Manager
./cs-whitelist-client --feature es.snapshot.frequent_snapshot_disabled_domains --target-stack swift-<region>-prod --target-client-id <aws-account-id> -r prod --region <region> --action list
# - moderate broken | disable hourly snapshots
aescloudopstools $REGION domain cp cs-whitelist-client-snapshot --feature
es.snapshot.frequent_snapshot_disabled_domains \
--args "--target-stack swift-$REGION-prod --target-client-id 755424219130 -r prod --region $REGION --feature_value --
action enable"
# - 2021/04/28 - Commands above for list and disable hourly snapshots is broken since python3
# - workaround is to mwinit --breakglass -o and connect in ssh into a Ressource Manager
```

```
./cs-whitelist-client --feature es.snapshot.frequent_snapshot_disabled_domains --target-stack swift-<region>-prod --target-client-id <aws-account-id> -r prod --region <region> --feature_value <domain-name> --action enable
# - moderate | disable daily snapshots
aescloudopstools $REGION generic run --cmd-name cs-override-es-domain-config --cmd-args "-D $DI --snapshot-hour -1"
# - moderate | disable daily snapshots
aescloudopstools $REGION generic run --cmd-name cs-override-es-domain-config --cmd-args "-D $DI --snapshot-hour XX"
# - moderate | delete stuck snapshot
```

```
tumbler $REGION aes domain dp node es curl delete -D $DI -u "_snapshot/cs-automated-er
```

```
# - high 2p | delete di
aescloudopstools run $REGION domain cp cs-delete-domain-instance --args "-D $DI -i <ident>"
```

DP Commands

_CLUSTER/SETTINGS COMMANDS

```
# - low | _cluster/settings all with defaults
tumbler $REGION aes domain dp node es curl get -D $DI -u '_cluster/settings?include_defaults=true&pretty'
# - low | check routing allocation require
tumbler $REGION aes domain dp node es curl get -D $DI -u '*/_settings?pretty' | \
grep -A 10 -B 10 "routing" | \
grep -A 10 -B 10 "allocation" | \
grep -A 10 -B 10 "require"
# - low | checking voting_config_exclusion
tumbler $REGION aes domain dp node es curl get -D $DI -u '_cluster/settings?include_defaults=true'
tumbler $REGION aes domain dp node es curl get -D $DI -u '_cluster/state?filter_path=metadata.cluster_coordination.voting_config_exclusions&pretty' \
--data '{"metadata": {"cluster_coordination": {"voting_config_exclusions": [ ]}}}'
# - moderate | speed shards alloc
tumbler $REGION aes domain dp node es curl put -D $DI -u "_cluster/settings?pretty" \
--data '{
"transient": {
"indices.recovery.max_concurrent_file_chunks": 5,
"cluster.routing.allocation.node_concurrent_recoveries": 5,
"indices.recovery.max_bytes_per_sec": "500mb"
}}'
# - moderate | _cluster/settings search.max_open_scroll_context
tumbler $REGION aes domain dp node es curl get -D $DI --url "_cluster/settings?pretty&include_defaults&flat_settings=true" \
| grep max_open_scroll_context
```

```
tumbler $REGION aes domain dp node es curl put -D $DI --url "/_cluster/settings?flat_settings=true" \
--data '{"persistent": {"search.max_open_scroll_context": 3000}, "transient": {"search.max_open_scroll_context": 3000}}'
# - moderate | _cluster/settings max_shards_per_node
tumbler $REGION aes domain dp node es curl put -D $DI -u '_cluster/settings' \
--data '{"transient": {"cluster.max_shards_per_node": 1010}}'
# - moderate | remove 7.x read-only
tumbler $REGION aes domain dp node es curl put -D $DI -u '_cluster/settings?pretty' \
--data '{"persistent": {"cluster.blocks.read_only": false}}'
# - high | es version < 7.x change minimum_master_nodes
tumbler $REGION aes domain dp node es curl put -D $DI -u '_cluster/settings?pretty' \
--data '{"transient": {"discovery.zen.minimum_master_nodes": 1}}'
# - high | es version > 7.x discovery.zen.minimum_master_nodes
tumbler $REGION aes domain dp node es curl post -D $DI -u '_cluster/voting_config_exclusions?node_ids=' --data '{}'
tumbler $REGION aes domain dp node es curl get -D $DI -u '_cluster/state?filter_path=metadata.cluster_coordination'
tumbler $REGION aes domain dp node es curl delete -D $DI -u '_cluster/voting_config_exclusions?wait_for_removal=false'
# - high | set routing allocation di exclusion (old command, permission revoked)
tumbler $REGION aes domain dp node es curl put -u '_cluster/settings?pretty' -D $DI \
--data '{"transient": {"cluster.routing.allocation.exclude.di_number": ""}}'
tumbler $REGION aes domain dp node es curl put -u '_cluster/settings?pretty' -D $DI \
--data '{"transient": {"cluster.routing.allocation.exclude.di_number": "7,6"}}'
# - high | set routing allocation di exclusion (new available command, need 2PR from Service Team)
tumbler <REGION> aes domain dp node cluster-routing-allocation --domain-identifier <DOMAIN-IDENTIFIER> --di <DI-
NUMBER> --instance-id <INSTANCE-ID> --approver-id <APPROVER-ID>
```

This wrapper is a command which is used to exclude a DI number in a cluster, it come with below **Guardrails**:

- Cluster has to be in Blue/Green deployment
- Total number of DI in the cluster should be at least 2
- Latest DI should not be excluded

Please note, due to existing bug, you need to provide a healthy instance id in this command so that you can run this command successfully. Service team engineer has awarded this bug and will fix this later.

- high | remove routing allocation zone exclusion

```
tumbler $REGION aes domain dp node es curl put -u '_cluster/settings?pretty' -D $DI \
--data '{"transient": {"cluster.routing.allocation.awareness.attributes": ""}}'
```

```
tumbler $REGION aes domain dp node es curl put -u '_cluster/settings?pretty' -D $DI \
--data '{"transient": {"cluster.routing.allocation.awareness.attributes": "zone"}}'
```

HEALTH & ALLOCATION COMMANDS

- low | _CLUSTER/HEALTH

TUMBLER \$REGION AES DOMAIN DP NODE ES CURL GET -D \$DI -U '_CLUSTER/HEALTH?PRETTY'

TUMBLER \$REGION AES DOMAIN DP NODE ES CURL GET -D \$DI -U '_CLUSTER/HEALTH?PRETTY&LOCAL=TRUE'

- low | _cluster/allocation

```
tumbler $REGION aes domain dp node es curl get -D $DI -u '_cluster/allocation/explain?pretty'
```

```
tumbler $REGION aes domain dp node es curl post -D $DI -u '_cluster/allocation/explain?pretty' \
```

```
--data '{"index": ".opendistro-ism-config", "shard": 0, "primary": true}'
```

- low | retry failed allocation

```
tumbler run $REGION aes domain dp node es curl cluster reroute-retry-failed -D $DI --submit
```

```
tumbler $REGION aes domain dp node es curl post -D $DI -u "/_cluster/reroute?retry_failed=true" --data '{}'
```

- **low** **change disk watermark settings**

```
tumbler $REGION aes domain dp node es curl put -D $DI --url "localhost:9200/_cluster/settings?" --data '{"persistent": {"cluster.routing.allocation.disk.watermark.low": "25gb", "cluster.routing.allocation.disk.watermark.high": "15gb", "cluster.routing.allocation.disk.watermark.flood_stage": "5gb"}}'
```

SNAPSHOT RELATED COMMANDS

- **LOW** **CHECK SNAPSHOT**

```
TUMBLER $REGION AES DOMAIN DP NODE ES CURL GET -D $DI -U '_CLUSTER/STATE?PRETTY' | GREP 'SNAPSHOT_DELETIONS' -A 5
TUMBLER $REGION AES DOMAIN DP NODE ES CURL GET -D $DI -U '_TASKS?PRETTY'
TUMBLER $REGION AES DOMAIN DP NODE ES CURL GET -D $DI -U '_SNAPSHOT/CS-AUTOMATED/_ALL?PRETTY&ERROR_TRACE=TRUECURL'
TUMBLER $REGION AES DOMAIN DP NODE ES CURL GET -D $DI -U '_CLUSTER/STATE?PRETTY' | GREP -A 1 '"STATE": "FAILED"' | GREP NODE | SORT | UNIQU -C
TUMBLER $REGION AES DOMAIN DP NODE ES CURL GET -D $DI -U '_CLUSTER/STATE?PRETTY' | GREP -A 1 '"STATE": "ABORTED"' | GREP NODE | SORT | UNIQU -C
```

- **moderate** **take / delete snapshot**

```
tumbler $REGION aes domain dp node es curl put -D $DI -u '_snapshot/cs-automated/' --data '{}'
```

```
tumbler $REGION aes domain dp node es curl delete $DI -u '_snapshot/cs-automated-enc/'
```

- **high** **list red indices + restore them from snapshots**

```
tumbler $REGION aes domain dp node es curl get -D $DI --url '_cat/indices?v&health=red'
```

```
aesccloudopstools run $REGION generic run --cmd-name="cs-manage-es-snapshots" --cmd-args="-D $DI --action restore --repository-name $RESPOSITORY --snapshot-name $LAST_SUCCESSFUL_SNAPSHOT --restore-indices $RED_INDICES"
```

- **high** **restore-index (with 2PR from Service Team)**

```
tumbler run <REGION> aes domain dp node restore-index --domain-identifier <DOMAIN-IDENTIFIER> --source-index-name <INDEX-NAME> --instance-id <INSTANCE-ID> --approver-id <APPROVER-ID>
```

This wrapper is used to delete an index from a node and restore the index from the latest successful snapshot, it come with below **Guardrails**:

- If there is node drop
- If the index is an internal index only
- If the index is a red index only
- If the index is present in the latest successful snapshot

Please note, due to existing bug, you need to provide a healthy instance id in this command so that you can run this command successfully. Service team engineer has awarded this bug and will fix this later.

Kibana Commands

- **moderate** **kibana fix index**

```
tumbler $REGION aes domain dp node es curl post -D $DI -u '_aliases' \
--data '{"actions": [{"remove": {"index": ".kibana_1", "alias": ".kibana"}}, {"add": {"index": ".kibana_2", "alias": ".kibana"}}]}'
```

```
tumbler $REGION aes domain dp node es curl post -D $DI -u '_reindex' \
--data '{"source": {"index": ".kibana_1"}, "dest": {"index": ".kibana_2"}}'
```

- **low** **check kibana api**

```
tumbler $REGION aes domain dp node shell -b "curl" -a "-s localhost:9200/_plugin/kibana/app/kibana" -D $DI
```

```
tumbler $REGION aes domain dp node shell -b "curl" -a "-s localhost:5601/api/status" -D $DI | jq -r ".status.overall"
```

BOUNCE COMMANDS

- HIGH TUMBLER BOUNCE ES

TUMBLER \$REGION AES DOMAIN DP NODE ES DEACTIVATE -D \$DI -i \$I
TUMBLER \$REGION AES DOMAIN DP NODE ES ACTIVATE -D \$DI -i \$I

- high tumbler bounce kibana

tumbler \$REGION aes domain dp node kibana deactivate -D \$DI -i \$I
tumbler \$REGION aes domain dp node kibana activate -D \$DI -i \$I

- high tumbler bounce sdpdagent

tumbler \$REGION aes domain dp node sdpdagent deactivate-agent -D \$DI —pip \$IP
tumbler \$REGION aes domain dp node sdpdagent activate-agent -D \$DI —pip \$IP

OTHER DP COMMANDS

- low shell find file

tumbler \$REGION aes domain dp node shell -D \$DI -b "find" -a " / -name *es_bootstrap*" -i i-0afe4672b284ac30d

- high 2p < 7.x seed master recovery

aescloudopstools run \$REGION domain cp cs-control-auto-scaling —args \
"-D \$DI —suspend —asg-number -i "
tumbler run \$REGION aes domain dp node haproxy deactivate -D \$DI -i \$I
tumbler \$REGION aes domain dp node shell -D \$DI -b "curl" -a \
"-s -XPOST localhost:9200/___aes/cluster/bootstrap -d '{\"initial_master_nodes\": [\"IP.IP.IP\"]}' -H \"Content-Type: application/json\" -i \$I/html
- - -

tumbler \$REGION aes domain dp node shell -D \$DI -b "curl" -a \
"-s -XPOST localhost:9200/___aes/shutdown -d '{\"StayDown\": \"false\"}' -H \"Content-Type: application/json\" -i \$I
OR

tumbler \$REGION aes domain dp node shell -b 'ps' -a ' -aux' -D \$DI -i \$I | grep 'bootstrap'
tumbler \$REGION aes domain dp node shell -b 'kill' -a " -D \$DI -i \$I
- - -

tumbler aes domain dp node es curl put -u "localhost:9200/_cluster/settings" -D —data '{"persistent":
{\"cluster.blocks.read_only\":\"true\"}}'
tumbler run \$REGION aes domain dp node haproxy activate -D \$DI -i \$I
aescloudopstools run \$REGION domain cp cs-control-auto-scaling —args "-D \$DI —resume —asg-number -i "
tumbler \$REGION aes domain dp node es curl get -D \$DI -u '_cluster/settings?include_defaults=true&&pretty&&flat_settings' |
grep "aes.jetty.cp_events.upgrade_status"

- low tail logs

tumbler \$REGION aes domain dp node shell -D \$DI -b "tail" —args "-n 1000 /apollo/env/swift-*prod-
Kibana/var/output/logs/kibana.log" -i \$I
tumbler \$REGION aes domain dp node shell -D \$DI -b "tail" —args "-n 1000 /apollo/env/swift-*prod-
Kibana/var/output/logs/kibana.log"

- low check pid of es process

tumbler \$REGION aes domain dp node shell -D \$DI -b 'ps' -a ' -wwef' -i \$I | grep -e "bootstrap"

- moderate check pid of cloudwatch and bounce

tumbler \$REGION aes domain dp node shell -b 'ps' -a ' -aux' -D \$DI -i \$I | grep CloudWatch | grep -v grep | awk '{print \$2}'
tumbler \$REGION aes domain dp node shell -b 'kill' -a " -D \$DI -i \$I

- moderate tumbler captures diagnostic dumps (light weight and doesn't take full heap dumps)

tumbler \$REGION aes domain dp node es profilers thread -D \$DI -i \$I && sleep 5 && \
tumbler \$REGION aes domain dp node es profilers thread -D \$DI -i \$I && \
tumbler \$REGION aes domain dp node diagnose -D \$DI -i \$I

#—moderate! increase the disk space— Avoid using this tool. Modifications made by this tool hinders the inplace updates on customer cluster. Recommend the customer to perform inplace volume update.

```
tumbler $REGION aes domain dp node modify-data-volume -D $DI -i --volume-size
```

- **low** grepping ES logs

```
tumbler $REGION aes domain dp node shell -D $DI -b "grep" -a " Exception /apollo/env/swift-*prod-*ES*/var/output/logs/elasticsearch.log" -i $I
```

- **high** reboot instance

```
tumbler --bypass $REGION aes domain dp node reboot-instance "-D $DI --instance-id $I"
```

- **moderate** closing index (permission revoked)

```
tumbler $REGION aes domain dp node es curl post -D $DI -u "/_close" --data '{}'
```

- **high** new closing index (need 2PR from Service Team)

```
tumbler run <REGION> aes domain dp node close-index --domain-identifier <DOMAIN-IDENTIFIER> --snapshot-id <SNAPSHOT-ID> --source-index-name <INDEX-NAME> --instance-id <INSTANCE-ID> --approver-id <APPROVER-ID>
```

This wrapper is used to close an index, it come with below **Guardrails**:

- The index is red
- Index must be present in the snapshot whose id has been provided

Please note, due to existing bug, you need to provide a healthy instance id in this command so that you can run this command successfully. Service team engineer has awarded this bug and will fix this later.

- **moderate** opening index

```
tumbler $REGION aes domain dp node es curl post -D $DI -u "/_open" --data '{}'
```

- **moderate** increase number of instances per domain

```
aescloudopstools run $REGION domain cp cs-update-limits --args "-D $DI --es-max-instance-count "
```

- **low** check ES version upgrade history

```
aescloudopstools $REGION domain cp es-print-upgrade-history --args "-D $DI"
```

- **low** check software version

```
aescloudopstools $REGION domain cp cs-list --args "-D $DI --instances -p" | grep "SoftwareVersion"
```

- **low** list VPC resources

```
aescloudopstools $REGION domain cp es-manage-vpc-interface --args "list-resources -D $DI"
```

- **low** disable admission control settings

```
tumbler $REGION aes domain dp node es curl put -D $DI --url "localhost:9200/_cluster/settings?" --data '{ "persistent" : { "aes.jetty.admission_control.global_cpu_usage.enable" : "false" }, "transient" : { "aes.jetty.admission_control.global_cpu_usage.enable" : "false" } }'
```

- **moderate** recover deleted domain

```
aescloudopstools $REGION domain cp cs-recover-deleted-domain --args "-D $DI"
```

- **moderate** change value of server.maxPayloadBytes (for domain version lower than OS 2.19)

STEP 1: Checking existing max payload

```
for i in tumbler $REGION aes domain dp get-domain-resources -D $DI | jq -r '.ec2Instances[]' | select(.nodeType=="dataNode") | .instanceId'
```

```
do{
```

```
echo ${i}
```

```
tumbler $REGION aes domain dp node shell -b su -a 'root -c "cat /apollo/env/swift-us-east-1-prod-OSD_2_9AMI-ES2/var/kibana/config/opensearch_dashboards.yml"' -D $DI -i $i | grep maxPayload
```

```
done
```

STEP 2: Deactivate kibana

```
yes | tumbler --bypass $REGION aes domain dp deactivate-kibana -D $DI
```

STEP 3: Changed max payload

```
for i in tumbler $REGION aes domain dp get-domain-resources -D $DI | jq -r '.ec2Instances[]' | select(.nodeType=="dataNode") | .instanceId'
```

```
do
```

```
echo $i
```

```
tumbler $REGION aes domain dp node shell -b "su" --args " root -c \"sed -i 's/server.maxPayloadBytes:
5242880/server.maxPayloadBytes: 31457280/' /apollo/env/swift-us-east-1-prod-OSD_2_9AMI-
ES2/var/kibana/config/opensearch_dashboards.yml\" -D $DI --instance-id $i
done
```

STEP 4: Activated kibana nodes

```
yes | tumbler --bypass $REGION aes domain dp activate-kibana -D $DI
```

STEP 5: Checked status of all nodes

```
tumbler $REGION aes domain dp get-domain-resources -D $DI | jq -r '.ec2Instances[] | select(.nodeType=="dataNode") |
.instanceId' | \ awk '{print $1}' | xargs -l {} sh -c 'echo {}';
tumbler $REGION aes domain dp node shell -b \"curl\" -a \"-s localhost:5601/api/status\" -D $DI -i {} | jq -r
\".status.overall.state\"
```

STEP 6: Verified all maxPayloadBytes were set to the requested settings

```
for i in tumbler $REGION aes domain dp get-domain-resources -D $DI | jq -r '.ec2Instances[] |
select(.nodeType=="dataNode") | .instanceId'
do
echo $i
tumbler $REGION aes domain dp node shell -b su -a 'root -c "cat /apollo/env/swift-us-east-1-prod-OSD_2_9AMI-
ES2/var/kibana/config/opensearch_dashboards.yml"' -D $DI -i $i | grep maxPayload
done
```

***Please note**, for domain version **OS 2.19 and higher**, we implement a customer requested feature which is persisting server.maxpayloadbytes in [opensearch_dashboards.yml](#) behind a feature flag for account level. After adding this feature flag, cx are able to request increase the server.maxpayloadbytes and the change will not be wiped out after B/G or node replacement. In this, to update **server.maxPayloadBytes for OS 2.19 and higher**, we need to follow [new SOP](#). However, SupOps engineers do not have permission for new SOP yet, if that is the case, then please contact Service Team Frontend oncall for help.

For more information:

<https://quip-amazon.com/O0zTAHCuF0XI>

<https://w.amazon.com/bin/view/Search/A9/Swift/Ops/Procedures/Kibana#Hserver.maxPayloadBytes28EnableFeatureFlag2Cst>
[opmanuallyupdatingconfig29](#)

- **moderate** | **restore domain from snapshot of deleted domain**

```
aescloudopstools run generic run --cmd-name es-tombstone-restore-snapshot --cmd-args "--deleted-domain-ident --
validate-deleted-domain --second-user= cp-on-call"
```

```
aescloudopstools run generic run --cmd-name es-tombstone-restore-snapshot --cmd-args "--deleted-domain-ident --new-
domain-ident --validate-new-domain --second-user= cp-on-call"
```

```
aescloudopstools run generic run --cmd-name es-tombstone-restore-snapshot --cmd-args "--deleted-domain-ident --new-
domain-ident --restore --second-user= cp-on-call"
```

```
aescloudopstools run generic run --cmd-name es-tombstone-restore-snapshot --cmd-args "--deleted-domain-ident --new-
domain-ident --delete-restore-repo --second-user= cp-on-call"
```

SOP:

<https://w.amazon.com/bin/view/Search/A9/Swift/Ops/Procedures/AccidentalDeleteRecovery/#HProcedureto restoresnapshotfro mdeleteddomain>